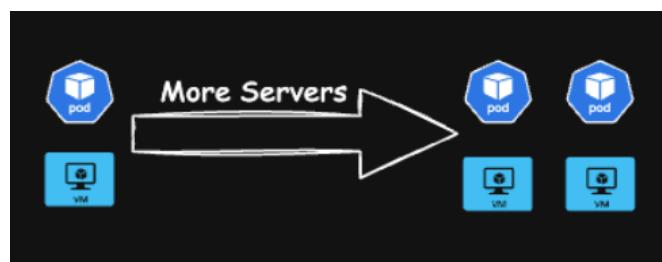


KUBERNETES HPA & VPA CONCEPTS

SCALING: Adjusting the resources based on varying load on application.

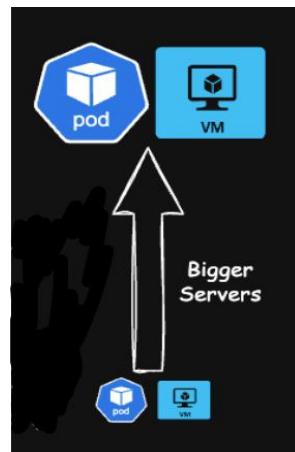
Pods can be scaled in following 2 ways in K8s:

- 1) Horizontal Scaling:** Increasing or Decreasing the number of pods or number of VMs based on application load.
 - **Manual Scaling:** By changing replicas count in deployment.yaml
 - **Automatic Scaling:** By using Horizontal pods Autoscaler



Horizontal Scaling

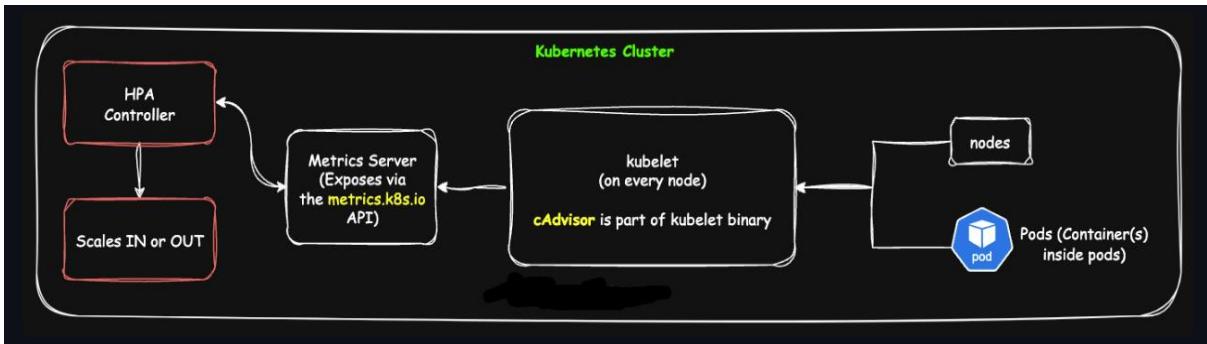
- 2) Vertical Scaling:** Increasing or decreasing the resources(CPU, Memory etc) of the same pod or VM based on traffic to application.
 - **Manual Scaling:** By changing container requests and limits in deployment.yaml
 - **Automatic Scaling:** By using Vertical pods Autoscaler



Vertical Scaling

- 1) Horizontal Pods Autoscaler(HPA):** It is a Kubernetes object that automatically **scales out(increase)** or **scales in(decrease)** the number of

pods based on Resource usage or Custom metrics. It runs as a control loop, checking metrics every 15 seconds.



Flow with Default Resource Metrics (CPU & Memory)

Note: Metric Server is the pre-requisite for HPA

Practical:

Deployment.yaml:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deploy
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
      resources:
        requests:
          cpu: "100m"
        limits:
          cpu: "200m"
```

```

  ports:
    - containerPort: 80

```

Service .yaml:

```

apiVersion: v1
kind: Service
metadata:
  name: nginx-svc
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: ClusterIP

```

Apply;

```

kubectl apply -f nginx-deployment.yaml
kubectl apply -f nginx-Service .yaml

```

Verify;

```

root@DESKTOP-C6P8EQS:~/kubernetes/11)HPA_VPA$ kubectl get all
NAME                               READY   STATUS    RESTARTS   AGE
pod/nginx-deploy-bb9f8c596-6s52p   1/1     Running   0          51s

NAME              TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/kubernetes   ClusterIP   10.96.0.1   <none>        443/TCP   4m40s
service/nginx-svc   ClusterIP   10.96.35.76  <none>        80/TCP    51s

NAME                  READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/nginx-deploy   1/1     1           1           52s

NAME                DESIRED   CURRENT   READY   AGE
replicaset.apps/nginx-deploy-bb9f8c596   1         1         1         51s

```

Lets create a HPA object:

```

kubectl autoscale deployment nginx-deploy --
cpu='50%' --min=1 --max=5

```

It states that at least 1 replica must run, and if CPU usage exceeds 50% of the request, scale up to a maximum of 5 replicas.

```

root@DESKTOP-C6P8EQS:~/kubernetes/11)HPA_VPA$ kubectl get hpa
NAME      REFERENCE      TARGETS      MINPODS   MAXPODS   REPLICAS   AGE
nginx-deploy  Deployment/nginx-deploy  cpu: 0%/50%  1          5          1          95s

```

Lets apply some CPU load on the pods using below command from 2 different tabs:

```
kubectl run -it --rm load-generator-1 --image=busybox -- /bin/sh -c "while true; do wget -q -O- http://nginx-svc; done"
```

Behavior:

- As load increases, CPU utilization rises.
- HPA will scale from 1 pod up to max 5 if needed.
- Once load decreases (stop load-generator), HPA scales back down.

CPU Utilization:

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
nginx-deploy	Deployment/nginx-deploy	cpu: 0%/50%	1	5	1	10m
nginx-deploy	Deployment/nginx-deploy	cpu: 3%/50%	1	5	1	11m
nginx-deploy	Deployment/nginx-deploy	cpu: 85%/50%	1	5	1	11m
nginx-deploy	Deployment/nginx-deploy	cpu: 95%/50%	1	5	2	11m
nginx-deploy	Deployment/nginx-deploy	cpu: 82%/50%	1	5	2	11m
nginx-deploy	Deployment/nginx-deploy	cpu: 80%/50%	1	5	2	12m

Scaling out of Replicas:

NAME	READY	STATUS	RESTARTS	AGE
nginx-deploy-bb9f8c596-6s52p	1/1	Running	0	12m
load-generator-1	0/1	Pending	0	0s
load-generator-1	0/1	Pending	0	0s
load-generator-1	0/1	ContainerCreating	0	0s
load-generator-1	1/1	Running	0	7s
nginx-deploy-bb9f8c596-6t2k8	0/1	Pending	0	0s
nginx-deploy-bb9f8c596-6t2k8	0/1	Pending	0	0s
nginx-deploy-bb9f8c596-6t2k8	0/1	ContainerCreating	0	0s
nginx-deploy-bb9f8c596-6t2k8	1/1	Running	0	28s
nginx-deploy-bb9f8c596-dggsz	0/1	Pending	0	0s
nginx-deploy-bb9f8c596-dggsz	0/1	Pending	0	0s
nginx-deploy-bb9f8c596-dggsz	0/1	ContainerCreating	0	0s
nginx-deploy-bb9f8c596-dggsz	1/1	Running	0	5s

Kubectl Events:

0s	Normal	Scheduled	Pod/load-generator-1	Successfully assigned default/load-generator-1 to rayeez-cluster-worker
0s	Normal	Pulling	Pod/load-generator-1	Pulling image "busybox"
0s	Normal	Pulled	Pod/load-generator-1	Successfully pulled image "busybox" in 7.568s (7.568s including waiting). Image size: 2224358 bytes.
0s	Normal	Created	Pod/load-generator-1	Created container load-generator-1
0s	Normal	Started	Pod/load-generator-1	Started container load-generator-1
0s	Normal	SuccessfulRescale	HorizontalPodAutoscaler/nginx-deploy	New size: 2; reason: cpu resource utilization (percentage of request) above target
0s	Normal	ScalingReplicaSet	Deployment/nginx-deploy	Scaled up replica set nginx-deploy-bb9f8c596 to 2 from 1
0s	Normal	SuccessfulCreate	ReplicaSet/nginx-deploy-bb9f8c596	Created pod: nginx-deploy-bb9f8c596-6t2k8
0s	Normal	Scheduled	Pod/nginx-deploy-bb9f8c596-6t2k8	Successfully assigned default/nginx-deploy-bb9f8c596-6t2k8 to rayeez-cluster-worker2
0s	Normal	Pulling	Pod/nginx-deploy-bb9f8c596-6t2k8	Pulling image "nginx"
0s	Normal	Pulled	Pod/nginx-deploy-bb9f8c596-6t2k8	Successfully pulled image "nginx" in 27.65s (27.65s including waiting). Image size: 59772801 bytes.
0s	Normal	Created	Pod/nginx-deploy-bb9f8c596-6t2k8	Created container nginx-container
0s	Normal	Started	Pod/nginx-deploy-bb9f8c596-6t2k8	Started container nginx-container
0s	Normal	SuccessfulRescale	HorizontalPodAutoscaler/nginx-deploy	New size: 3; reason: cpu resource utilization (percentage of request) above target
0s	Normal	ScalingReplicaSet	Deployment/nginx-deploy	Scaled up replica set nginx-deploy-bb9f8c596 to 3 from 2
0s	Normal	SuccessfulCreate	ReplicaSet/nginx-deploy-bb9f8c596	Created pod: nginx-deploy-bb9f8c596-dggsz
0s	Normal	Scheduled	Pod/nginx-deploy-bb9f8c596-dggsz	Successfully assigned default/nginx-deploy-bb9f8c596-dggsz to rayeez-cluster-worker
0s	Normal	Pulling	Pod/nginx-deploy-bb9f8c596-dggsz	Pulling image "nginx"
0s	Normal	Pulled	Pod/nginx-deploy-bb9f8c596-dggsz	Successfully pulled image "nginx" in 2.654s (2.654s including waiting). Image size: 59772801 bytes.
0s	Normal	Created	Pod/nginx-deploy-bb9f8c596-dggsz	Created container nginx-container
0s	Normal	Started	Pod/nginx-deploy-bb9f8c596-dggsz	Started container nginx-container

Applied Load generator-2;

```
kubectl run -it --rm load-generator-2 --image=busybox -- /bin/sh -c "while true; do wget -q -O- http://nginx-svc; done"
```

CPU Utilization:

nginx-deploy	Deployment/nginx-deploy	cpu: 31%/50%	1	5	3	20m
nginx-deploy	Deployment/nginx-deploy	cpu: 36%/50%	1	5	3	20m
nginx-deploy	Deployment/nginx-deploy	cpu: 30%/50%	1	5	3	20m
nginx-deploy	Deployment/nginx-deploy	cpu: 38%/50%	1	5	3	21m
nginx-deploy	Deployment/nginx-deploy	cpu: 30%/50%	1	5	3	21m
nginx-deploy	Deployment/nginx-deploy	cpu: 65%/50%	1	5	3	21m

Scaling out of Replicas:

load-generator-2	0/1	Pending	0	0s
load-generator-2	0/1	Pending	0	0s
load-generator-2	0/1	ContainerCreating	0	0s
load-generator-2	1/1	Running	0	10s
nginx-deploy-bb9f8c596-d425x	0/1	Pending	0	0s
nginx-deploy-bb9f8c596-d425x	0/1	Pending	0	0s
nginx-deploy-bb9f8c596-d425x	0/1	ContainerCreating	0	0s
nginx-deploy-bb9f8c596-d425x	1/1	Running	0	4s
nginx-deploy-bb9f8c596-9g595	0/1	Pending	0	0s
nginx-deploy-bb9f8c596-9g595	0/1	Pending	0	0s
nginx-deploy-bb9f8c596-9g595	0/1	ContainerCreating	0	0s
nginx-deploy-bb9f8c596-9g595	1/1	Running	0	5s

Kubectl Events:

0s	Normal	Scheduled	Pod/load-generator-2	Successfully assigned default/load-generator-2 to rayeez-cluster-worker2
0s	Normal	Pulling	Pod/load-generator-2	Pulling image "busybox"
0s	Normal	Pulled	Pod/load-generator-2	Successfully pulled image "busybox" in 7.49s (7.49s including waiting). Image size: 2224358 bytes.
0s	Normal	Created	Pod/load-generator-2	Created container load-generator-2
0s	Normal	Started	Pod/load-generator-2	Started container load-generator-2
0s	Normal	SuccessfulRescale	HorizontalPodAutoscaler/nginx-deploy	New size: 4; reason: cpu resource utilization (percentage of request) above target
0s	Normal	ScalingReplicaSet	Deployment/nginx-deploy-bb9f8c596	Scaled up replica set nginx-deploy-bb9f8c596 to 4 from 3
0s	Normal	SuccessfulCreate	ReplicaSet/nginx-deploy-bb9f8c596	Scaled pod nginx-deploy-bb9f8c596-d425x
0s	Normal	Scheduled	Pod/nginx-deploy-bb9f8c596-d425x	Successfully assigned default/nginx-deploy-bb9f8c596-d425x to rayeez-cluster-worker2
0s	Normal	Pulling	Pod/nginx-deploy-bb9f8c596-d425x	Pulling image "nginx"
0s	Normal	Pulled	Pod/nginx-deploy-bb9f8c596-d425x	Successfully pulled image "nginx" in 2.317s (2.317s including waiting). Image size: 59772801 bytes.
0s	Normal	Created	Pod/nginx-deploy-bb9f8c596-d425x	Created container nginx-container
0s	Normal	Started	Pod/nginx-deploy-bb9f8c596-d425x	Started container nginx-container
0s	Normal	SuccessfulRescale	HorizontalPodAutoscaler/nginx-deploy	New size: 5; reason: cpu resource utilization (percentage of request) above target
0s	Normal	ScalingReplicaSet	Deployment/nginx-deploy-bb9f8c596	Scaled up replica set nginx-deploy-bb9f8c596 to 5 from 4
0s	Normal	SuccessfulCreate	ReplicaSet/nginx-deploy-bb9f8c596	Created pod nginx-deploy-bb9f8c596-9g595
0s	Normal	Scheduled	Pod/nginx-deploy-bb9f8c596-9g595	Successfully assigned default/nginx-deploy-bb9f8c596-9g595 to rayeez-cluster-worker
0s	Normal	Pulling	Pod/nginx-deploy-bb9f8c596-9g595	Pulling image "nginx"
0s	Normal	Pulled	Pod/nginx-deploy-bb9f8c596-9g595	Successfully pulled image "nginx" in 2.71s (2.71s including waiting). Image size: 59772801 bytes.
0s	Normal	Created	Pod/nginx-deploy-bb9f8c596-9g595	Created container nginx-container
0s	Normal	Started	Pod/nginx-deploy-bb9f8c596-9g595	Started container nginx-container

Kubectl describe hpa nginx-deploy

root@DESKTOP-C6P8EQS:/# kubectl describe hpa nginx-deploy
Name: default
Namespace: default
Labels: <none>
Annotations: <none>
CreationTimestamp: Sat, 22 Nov 2025 10:30:22 +0000
Reference: Deployment/nginx-deploy
Metrics:
resource cpu on pods (as a percentage of request): 35% (35m) / 50%
Min replicas: 1
Max replicas: 5
Deployment pods: 5 current / 5 desired
Conditions:
Type Status Reason Message
AbleToScale True ScaleDownStabilized recent recommendations were higher than current one, applying the highest recent recommendation
ScalingActive True ValidMetricFound the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
ScalingLimited False DesiredWithinRange the desired count is within the acceptable range
Events:
Type Reason Age From Message
Normal SuccessfulRescale 23m horizontal-pod-autoscaler New size: 2; reason: cpu resource utilization (percentage of request) above target
Normal SuccessfulRescale 22m horizontal-pod-autoscaler New size: 3; reason: cpu resource utilization (percentage of request) above target
Normal SuccessfulRescale 13m horizontal-pod-autoscaler New size: 4; reason: cpu resource utilization (percentage of request) above target
Normal SuccessfulRescale 4m45s horizontal-pod-autoscaler New size: 5; reason: cpu resource utilization (percentage of request) above target

Lets Terminate Load generator-2 and observe scaling in behaviour;

load-generator-2	1/1	Terminating	0	18m
load-generator-2	0/1	Error	0	18m
load-generator-2	0/1	Error	0	18m
load-generator-2	0/1	Error	0	18m

CPU Utilization:

nginx-deploy	Deployment/nginx-deploy	cpu: 20%/50%	1	5	5	43m
nginx-deploy	Deployment/nginx-deploy	cpu: 19%/50%	1	5	5	43m
nginx-deploy	Deployment/nginx-deploy	cpu: 25%/50%	1	5	5	43m
nginx-deploy	Deployment/nginx-deploy	cpu: 18%/50%	1	5	5	43m

Scaling in of Replicas:

nginx-deploy-bb9f8c596-9g595	1/1	Terminating	0	13m
nginx-deploy-bb9f8c596-9g595	0/1	Completed	0	13m
nginx-deploy-bb9f8c596-9g595	0/1	Completed	0	13m
nginx-deploy-bb9f8c596-9g595	0/1	Completed	0	13m
nginx-deploy-bb9f8c596-dggsz	1/1	Terminating	0	31m
nginx-deploy-bb9f8c596-dggsz	0/1	Completed	0	31m
nginx-deploy-bb9f8c596-dggsz	0/1	Completed	0	31m
nginx-deploy-bb9f8c596-dggsz	0/1	Completed	0	31m

Kubectl Events:

0s	Normal	SuccessfulRescale	HorizontalPodAutoscaler/nginx-deploy	New size: 4; reason: All metrics below target
0s	Normal	ScalingReplicaSet	Deployment/nginx-deploy	Scaled down replica set nginx-deploy-bb9f8c596 to 4 from 5
0s	Normal	SuccessfulDelete	ReplicaSet/nginx-deploy-bb9f8c596	Deleted pod: nginx-deploy-bb9f8c596-9g595
0s	Normal	Killing	Pod/nginx-deploy-bb9f8c596-9g595	Stopping container nginx-container
0s	Normal	SuccessfulRescale	HorizontalPodAutoscaler/nginx-deploy	New size: 3; reason: All metrics below target
0s	Normal	ScalingReplicaSet	Deployment/nginx-deploy	Scaled down replica set nginx-deploy-bb9f8c596 to 3 from 4
0s	Normal	SuccessfulDelete	ReplicaSet/nginx-deploy-bb9f8c596	Deleted pod: nginx-deploy-bb9f8c596-dggsz
0s	Normal	Killing	Pod/nginx-deploy-bb9f8c596-dggsz	Stopping container nginx-container

Conclusion:

HPA automatically scales out and scales in the number of replicas on the cluster based on load/traffic received by the application.

2) Vertical Pods Autoscaler(VPA):

It automatically adjust the **CPU** and **memory requests** and **limits** of pods based on their actual usage.

- It increases or decreases the CPU & memory reservations of pods.
- Pods always have enough resources to perform efficiently.
- Wasted resources are minimized, avoiding over-provisioning.
- Removes guesswork. No need to manually tune CPU/memory requests.

Prerequisite: VPA recommender, Updater, Admission controller and metric Server.

Practical:

deployment.yaml:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deploy
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          resources:
            requests:
              cpu: "100m"
            limits:
              cpu: "200m"
          ports:
            - containerPort: 80
```

Service.yaml:

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-svc
spec:
  selector:
```

```

  app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
    type: ClusterIP

```

vpa.yaml:

```

apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: nginx-vpa
spec:
  targetRef:
    apiVersion: "apps/v1"
    kind: Deployment
    name: nginx-deploy
  updatePolicy:
    # updateMode options:
    # "Off" - VPA only recommends resources, does NOT apply them.
    # "Initial" - VPA sets recommended resources at pod creation, no changes after.
    # "Auto" - VPA automatically updates resources and restarts pods as needed.
    updateMode: "Auto"

```

Apply;

```

kubectl apply -f nginx-deployment.yaml
kubectl apply -f nginx-Service .yaml

```

Verify;

```

root@DESKTOP-C6P8EQS:~/kubernetes/11)HPA_VPA$ kubectl get all
NAME                           READY   STATUS    RESTARTS   AGE
pod/nginx-deploy-bb9f8c596-86hwq   1/1     Running   0          43s
pod/nginx-deploy-bb9f8c596-nlccm   1/1     Running   0          43s

NAME              TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes   ClusterIP   10.96.0.1      <none>           443/TCP     91m
service/nginx-svc   ClusterIP   10.96.104.63   <none>           80/TCP      43s

NAME                  READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/nginx-deploy   2/2     2           2           44s

NAME            DESIRED   CURRENT   READY   AGE
replicaset.apps/nginx-deploy-bb9f8c596   2         2         2         43s

```

We have set this CPU request and limit in yaml manifest:

```
Limits:  
  cpu: 200m  
Requests:  
  cpu: 100m
```

Lets create VPA and observe its behaviour;

```
kubectl apply -f vpa.yaml
```

Verify;

```
root@DESKTOP-C6P8EQS:~/kubernetes/11)HPA_VPA$ kubectl get vpa  
NAME      MODE   CPU    MEM    PROVIDED   AGE  
nginx-vpa Auto   25m   250Mi  True       37s
```

The moment we have created VPA; This event took place:

```
0s   Normal  EvictedPod  VerticalPodAutoscaler/nginx-vpa  VPA Updater evicted Pod nginx-deploy-bb9f8c596-bdcrq to apply resource recommendation.  
0s   Normal  SuccessfulCreate  ReplicaSet/nginx-deploy-bb9f8c596  Created pod: nginx-deploy-bb9f8c596-gnwkf  
0s   Normal  Scheduled  Pod/nginx-deploy-bb9f8c596-gnwkf  Successfully assigned default/nginx-deploy-bb9f8c596-gnwkf to rayeez-cluster-worker2  
0s   Normal  Pulling    Pod/nginx-deploy-bb9f8c596-gnwkf  Pulling image "nginx"  
0s   Normal  Pulled     Pod/nginx-deploy-bb9f8c596-gnwkf  Successfully pulled image "nginx" in 2.224s (2.224s including waiting). Image size: 597728  
01 bytes.  
0s   Normal  Created    Pod/nginx-deploy-bb9f8c596-gnwkf  Created container nginx-container  
0s   Normal  Started    Pod/nginx-deploy-bb9f8c596-gnwkf  Started container nginx-container
```

Now the CPU requests and limits are adjusted to following values by VPA.

```
Limits:  
  cpu: 50m  
Requests:  
  cpu: 25m  
  memory: 250Mi
```

Lets apply some CPU load on the pods using below command from 2 different tabs:

```
kubectl run -it --rm load-generator-1 --  
image=busybox -- /bin/sh -c "while true; do  
wget -q -O- http://nginx-svc; done"
```

Now Because of increasing CPU load on pods, VPA has set new CPU request, limits values to pods:

nginx-vpa	Auto	25m	250Mi	True	6m21s
nginx-vpa	Auto	25m	250Mi	True	7m19s
nginx-vpa	Auto	25m	250Mi	True	8m18s
nginx-vpa	Auto	25m	250Mi	True	9m14s
nginx-vpa	Auto	49m	250Mi	True	10m
nginx-vpa	Auto	49m	250Mi	True	11m
nginx-vpa	Auto	63m	250Mi	True	12m
nginx-vpa	Auto	63m	250Mi	True	12m
nginx-vpa	Auto	63m	250Mi	True	13m

Kubectl Events:

0s	Normal	Killing	Pod/nginx-deploy-bb9f8c596-f8f9f	Stopping container nginx-container
0s	Normal	EvictedByVPA	Pod/nginx-deploy-bb9f8c596-f8f9f	Pod was evicted by VPA Updater to apply resource recommendation.
0s	Normal	EvictedPod	VerticalPodAutoscaler/nginx-vpa	VPA Updater evicted Pod nginx-deploy-bb9f8c596-f8f9f to apply resource recommendation.
0s	Normal	SuccessfulCreate	ReplicaSet/nginx-deploy-bb9f8c596	Created pod: nginx-deploy-bb9f8c596-f8f9f
0s	Normal	Scheduled	Pod/nginx-deploy-bb9f8c596-f8f9f	Successfully assigned default/nginx-deploy-bb9f8c596-f8f9f to rayeez-cluster-worker
0s	Normal	Pulling	Pod/nginx-deploy-bb9f8c596-f8f9f	Pulling image "nginx"
0s	Normal	Pulled	Pod/nginx-deploy-bb9f8c596-f8f9f	Successfully pulled image "nginx" in 2.34s (2.34s including waiting). Image size: 59772801 bytes.
0s	Normal	Created	Pod/nginx-deploy-bb9f8c596-f8f9f	Created container nginx-container
0s	Normal	Started	Pod/nginx-deploy-bb9f8c596-f8f9f	Started container nginx-container

The VPA evicted one of the running pods, updated the CPU requests and limits, and then the ReplicaSet controller created new pods with these updated resource values.

Kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
load-generator-1	1/1	Running	0	30m
load-generator-2	1/1	Running	0	18m
nginx-deploy-bb9f8c596-f8f9f	1/1	Running	0	14m
nginx-deploy-bb9f8c596-pmlf7	1/1	Running	0	13m

Now the CPU request, Limits are set as below in each pods:

```
Limits:  
  cpu: 126m  
Requests:  
  cpu:       63m  
  memory:   250Mi
```

Conclusion: VPA dynamically adjusts the pod's resource requests and limits based on the application's varying load, while keeping the number of pods unchanged.